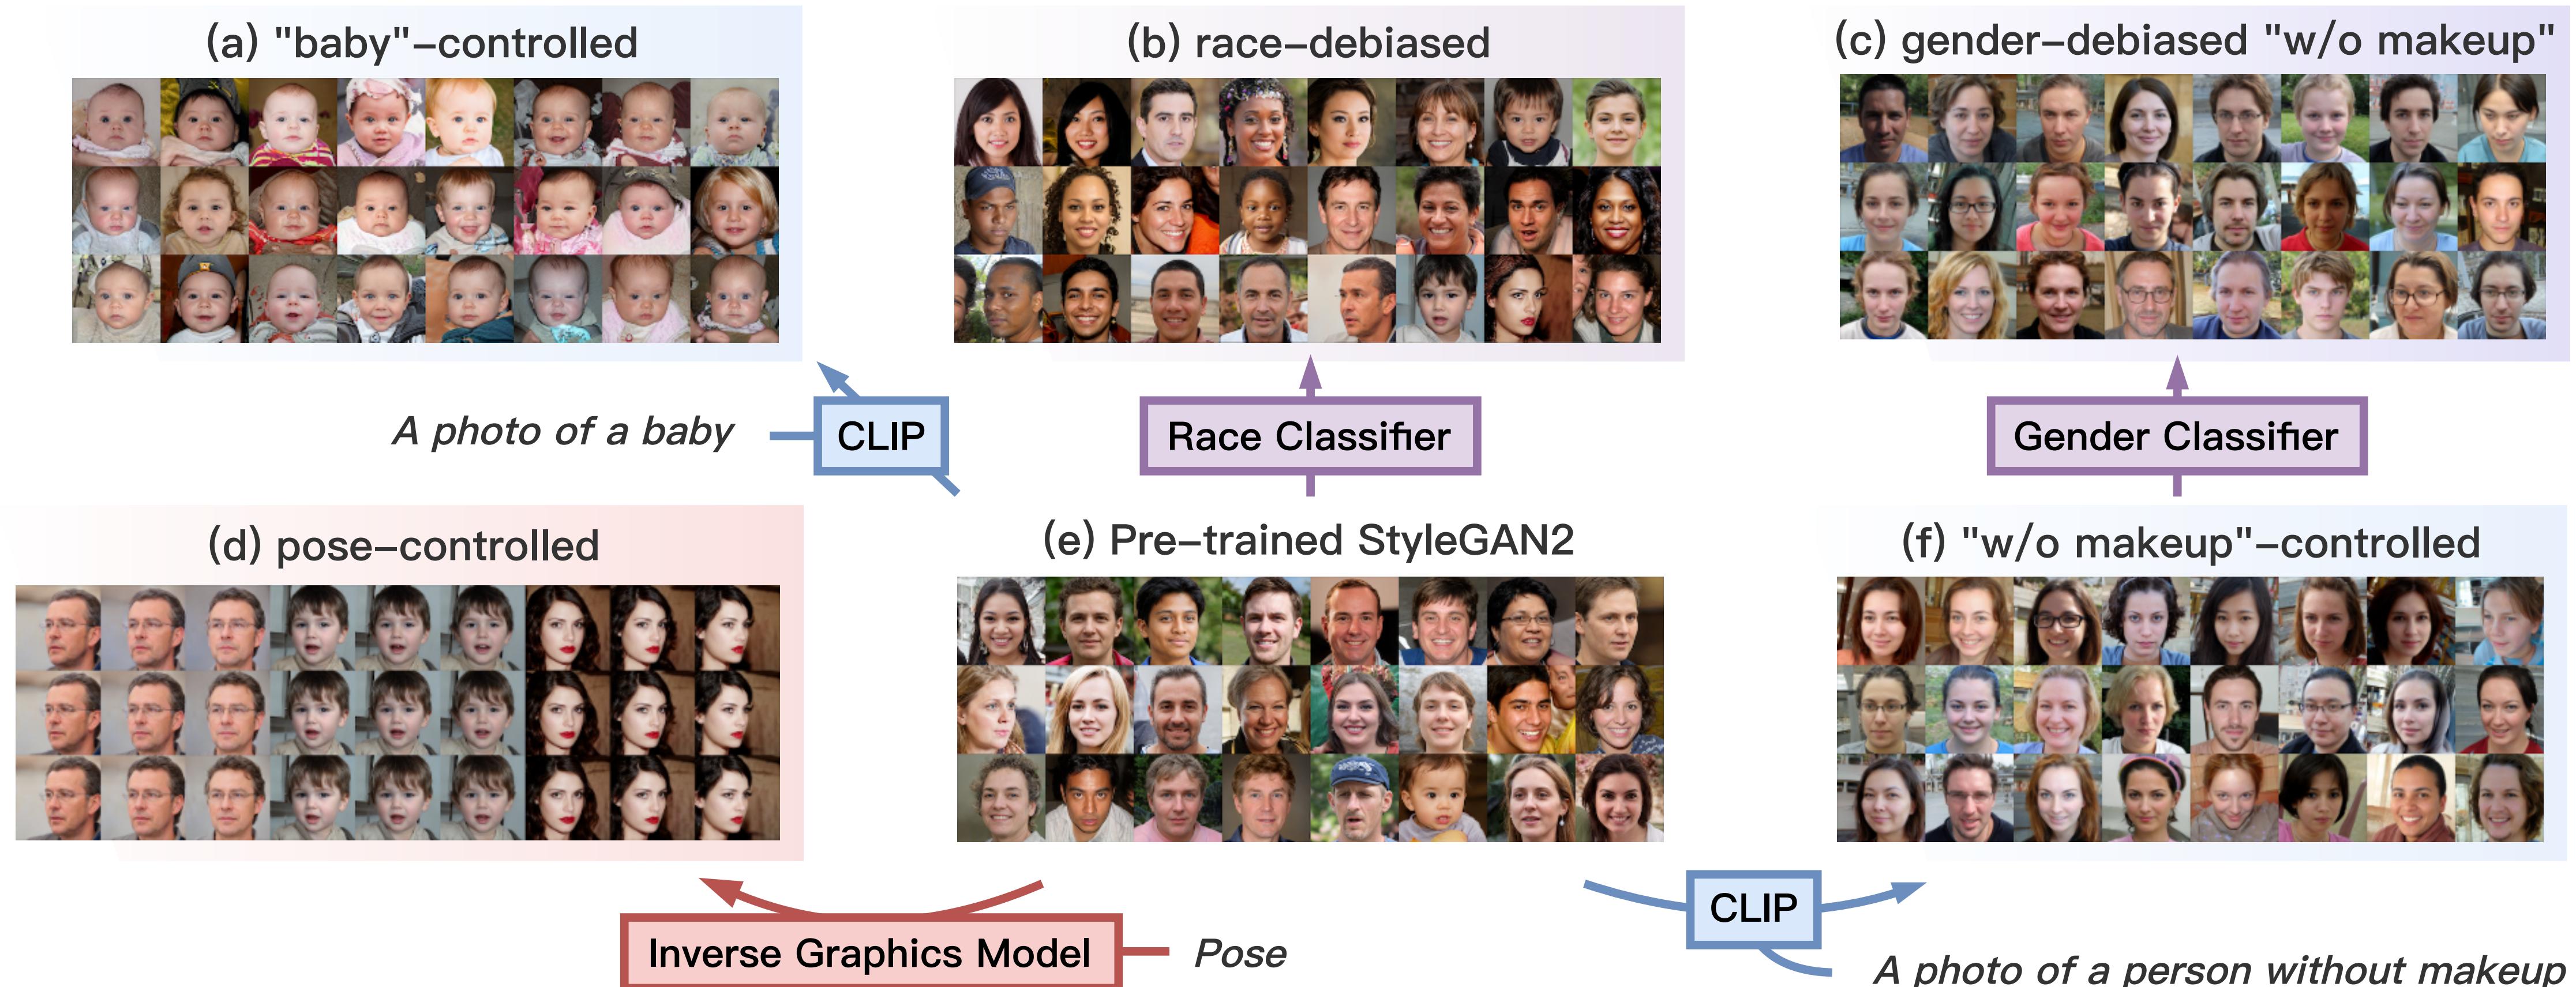
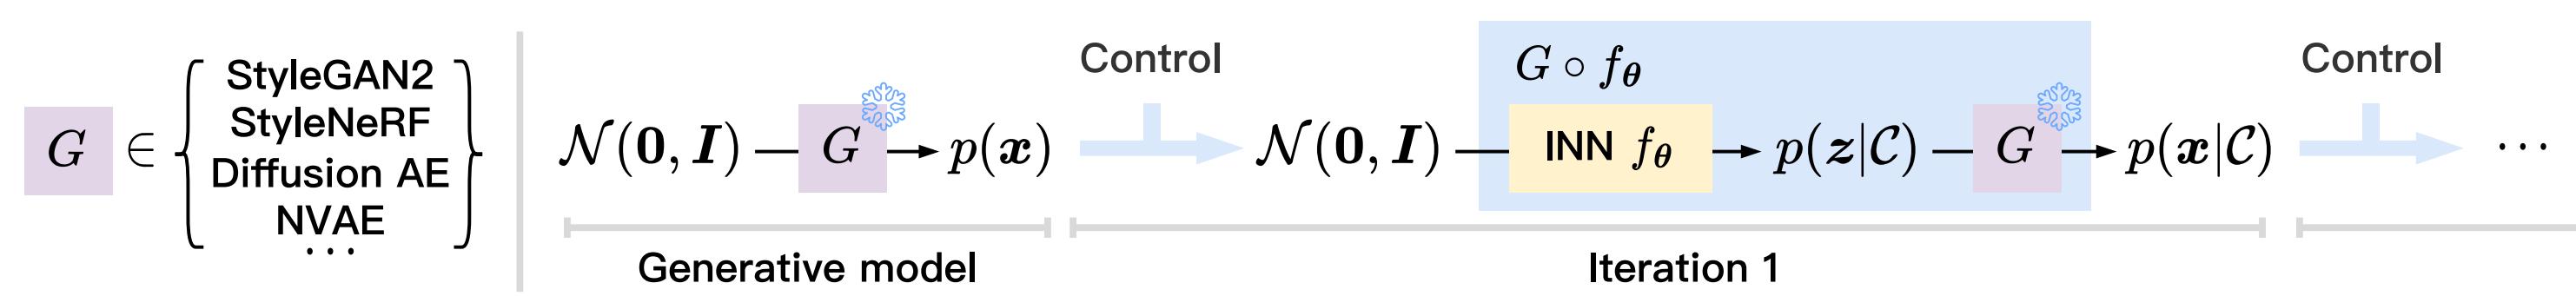


Generative Visual Prompt: Unifying Distributional Control of Pre-Trained Generative Models

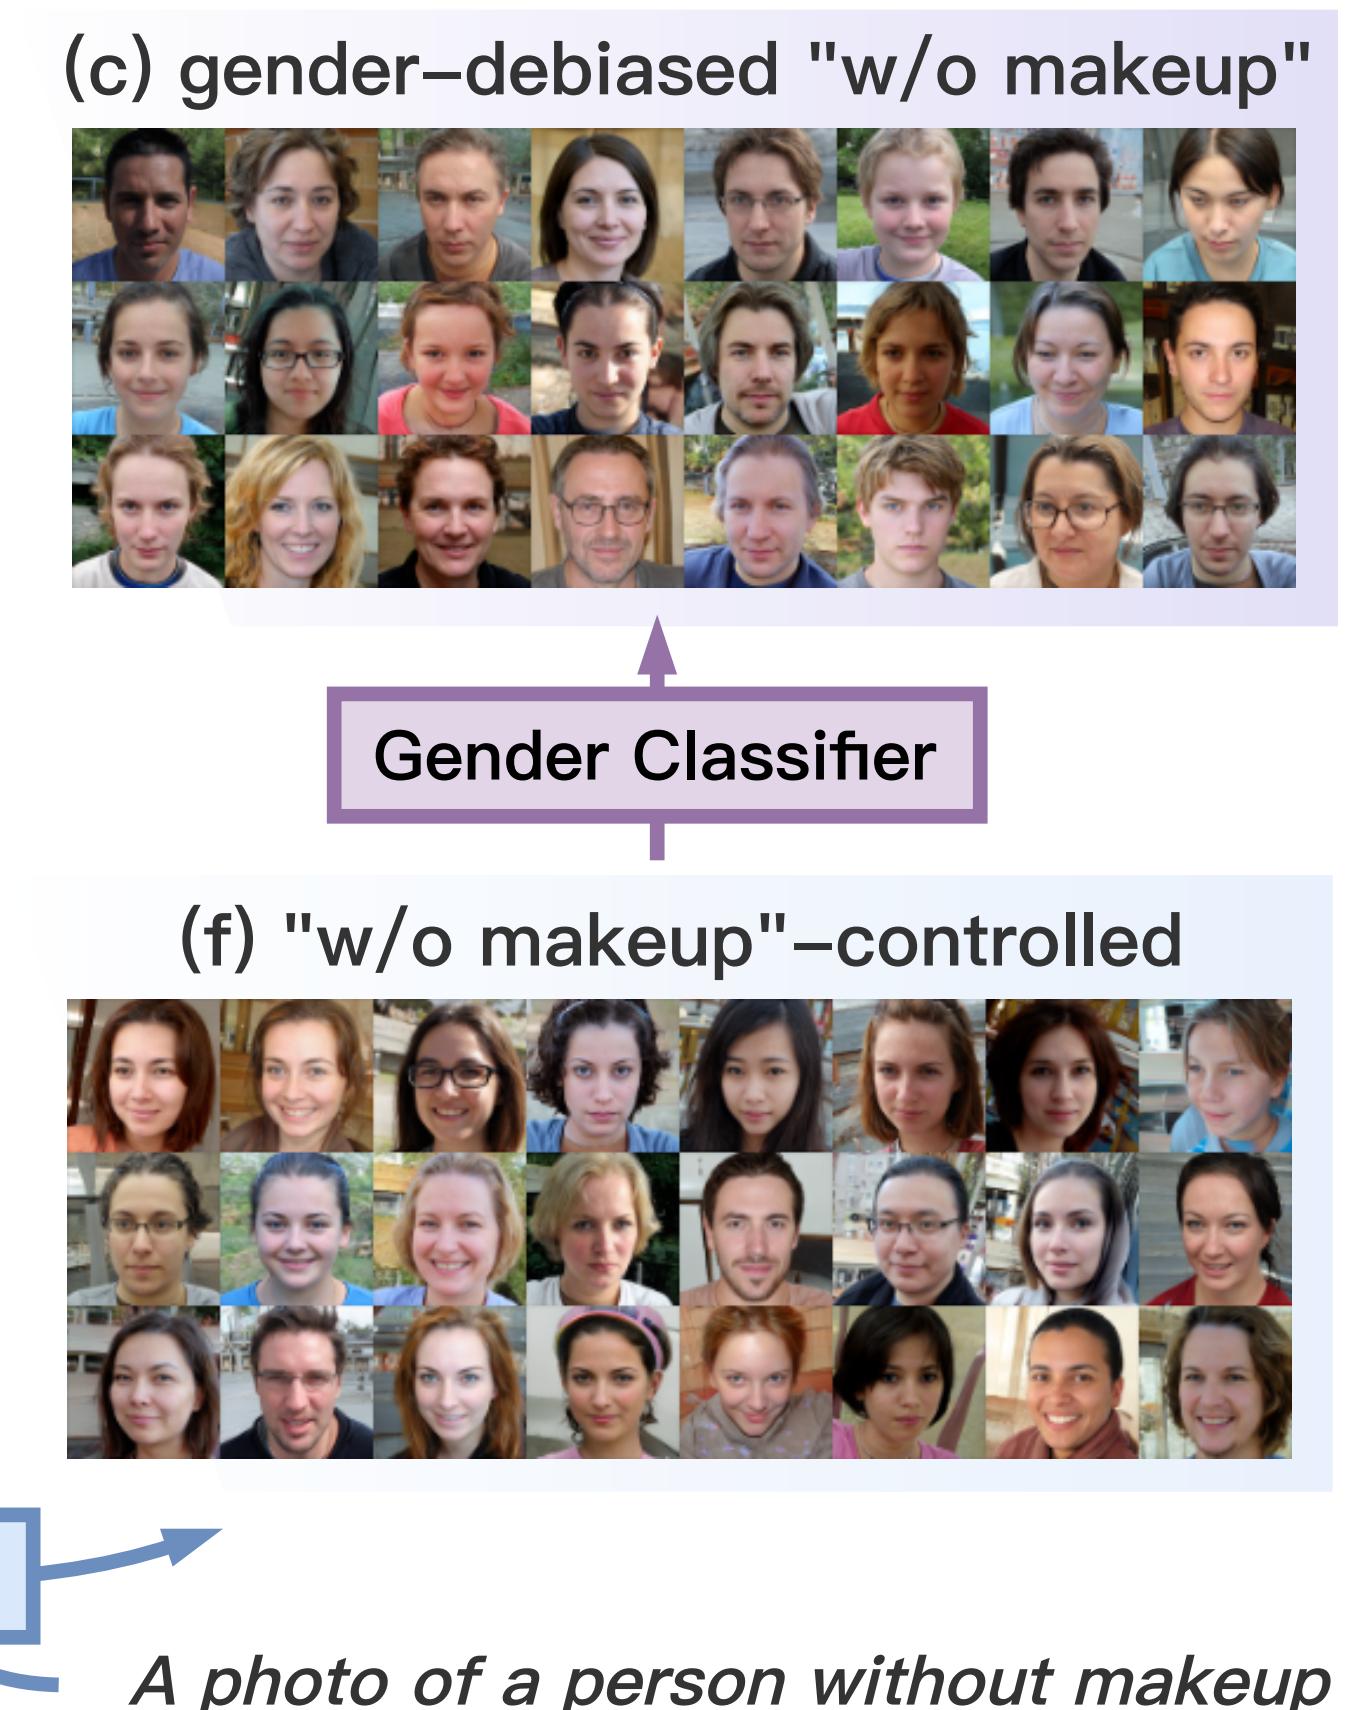
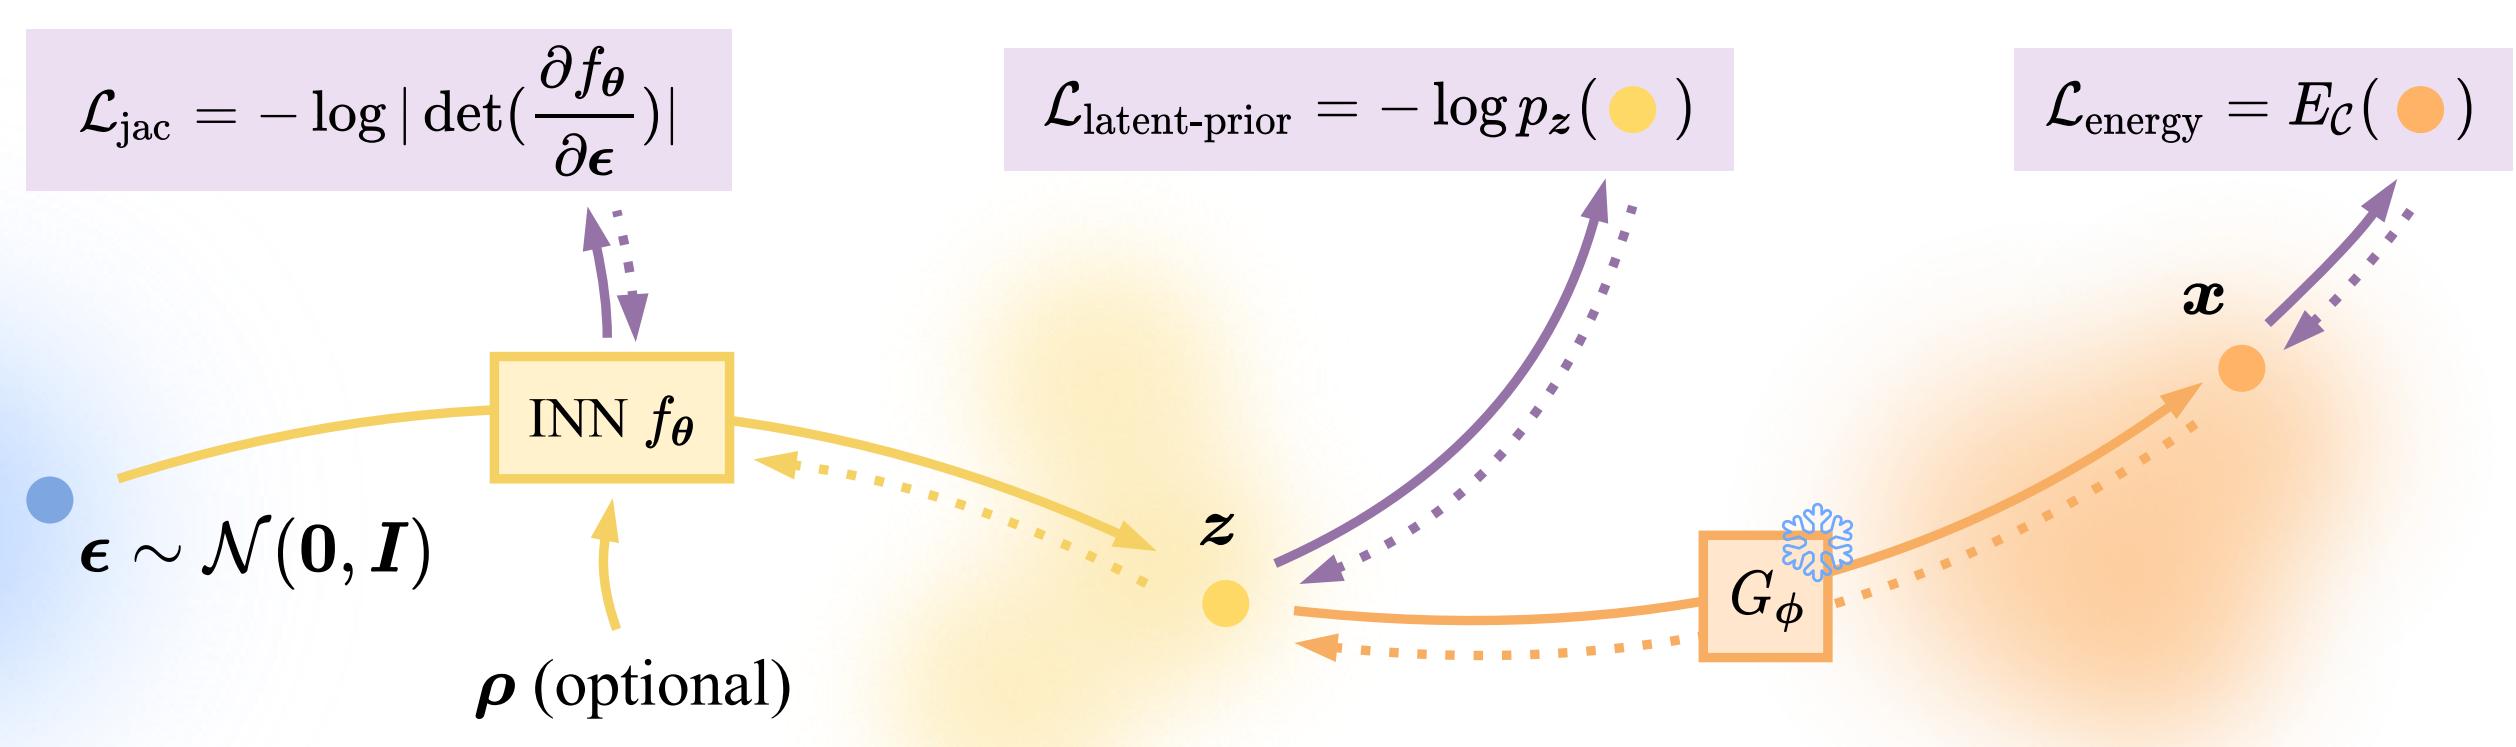
Chen Henry Wu, Saman Motamed, Shaunk Srivastava, Fernando De la Torre
Carnegie Mellon University



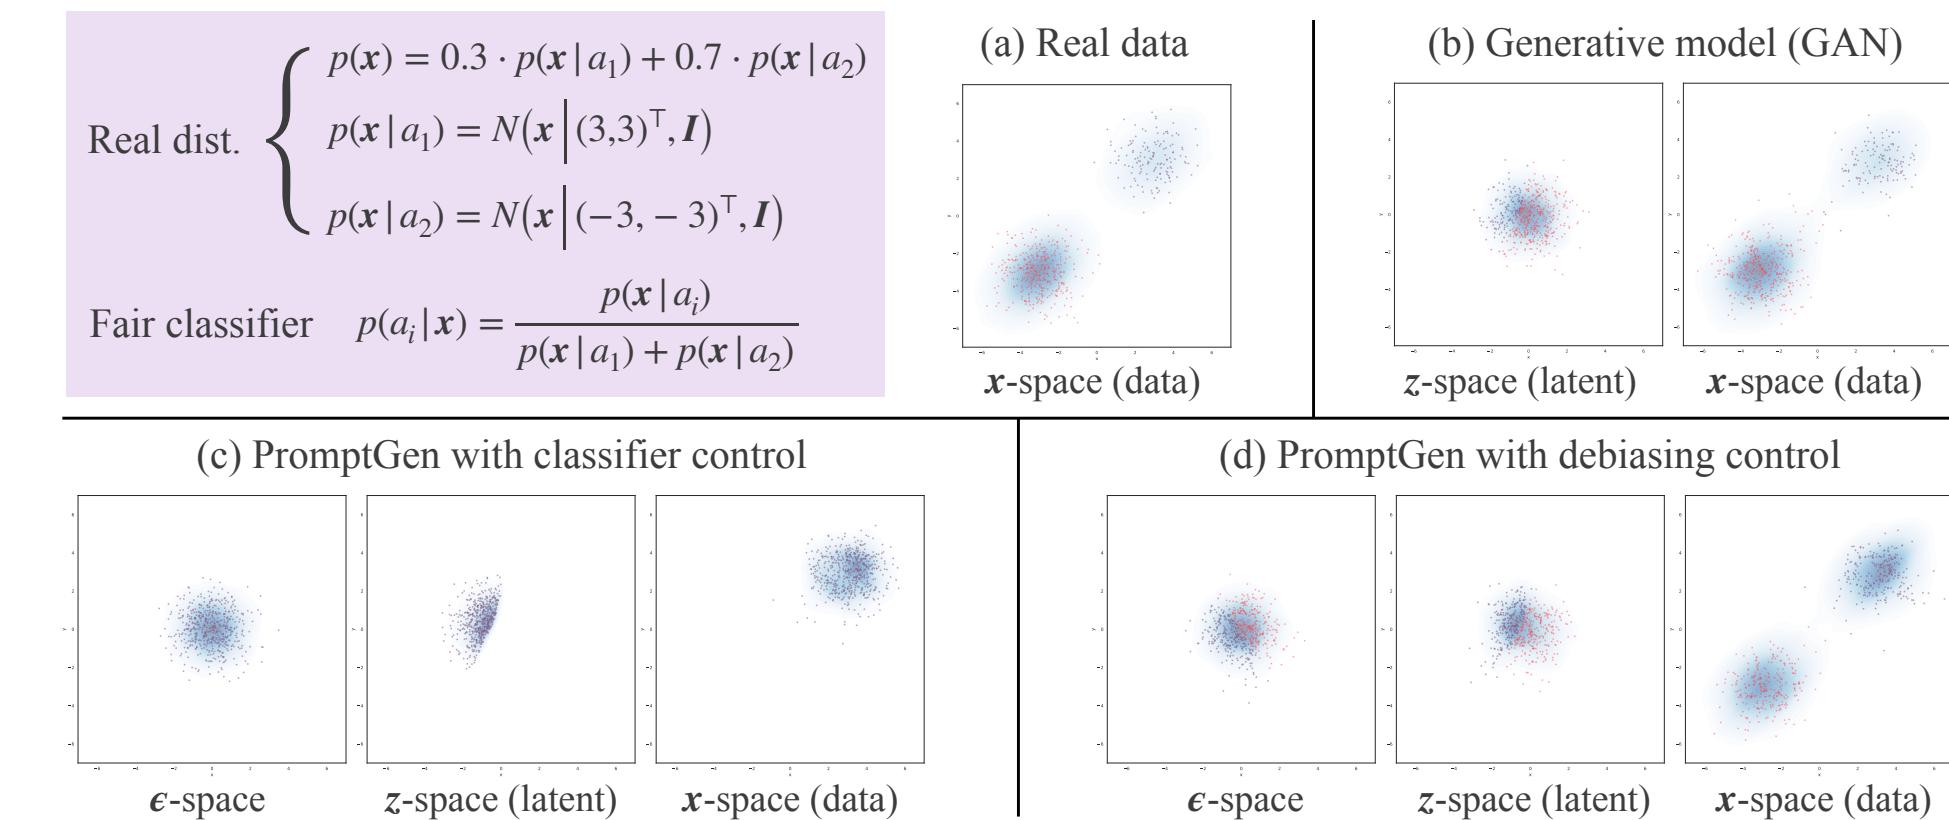
PromptGen Framework



PromptGen Training

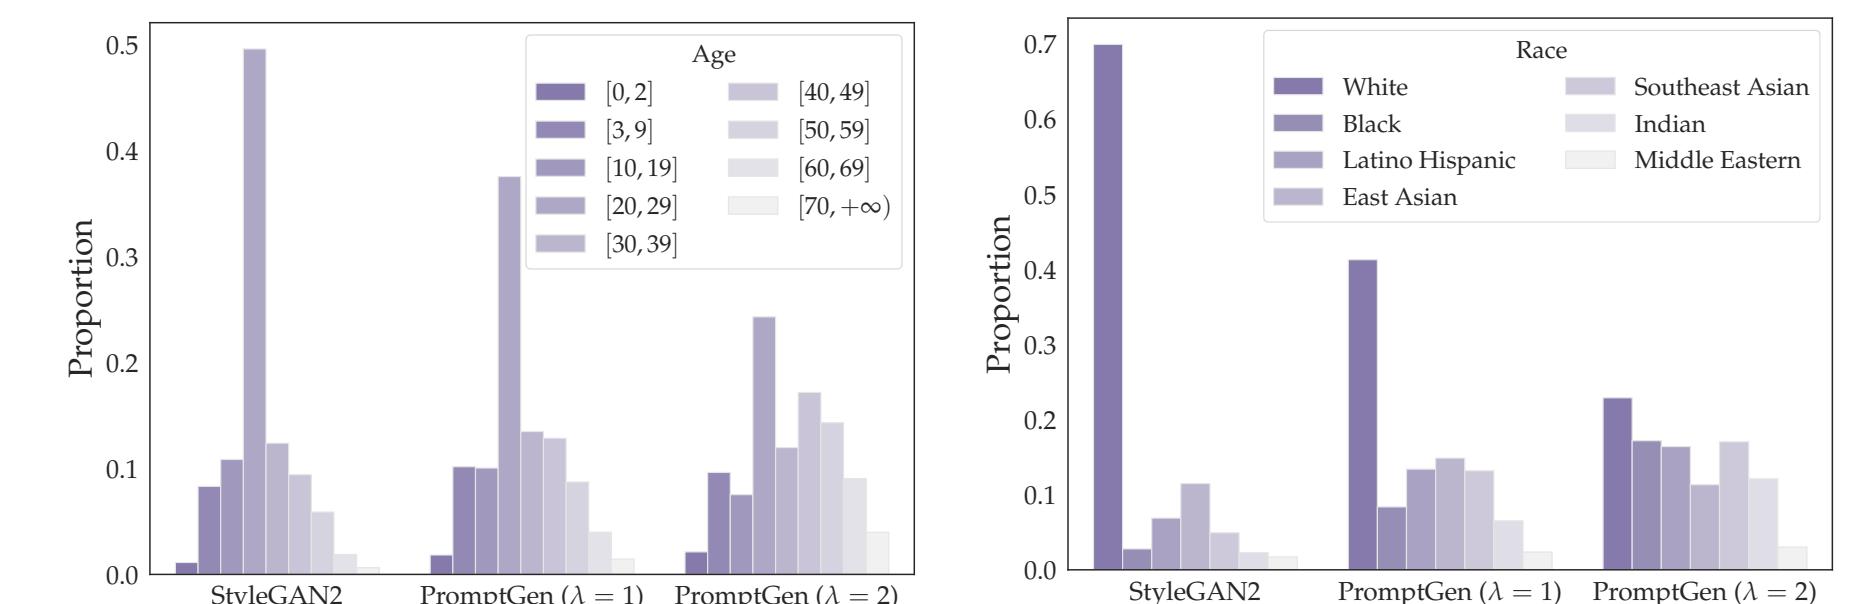


Synthetic Example



More Results

De-biasing StyleGAN2 with moment constraint



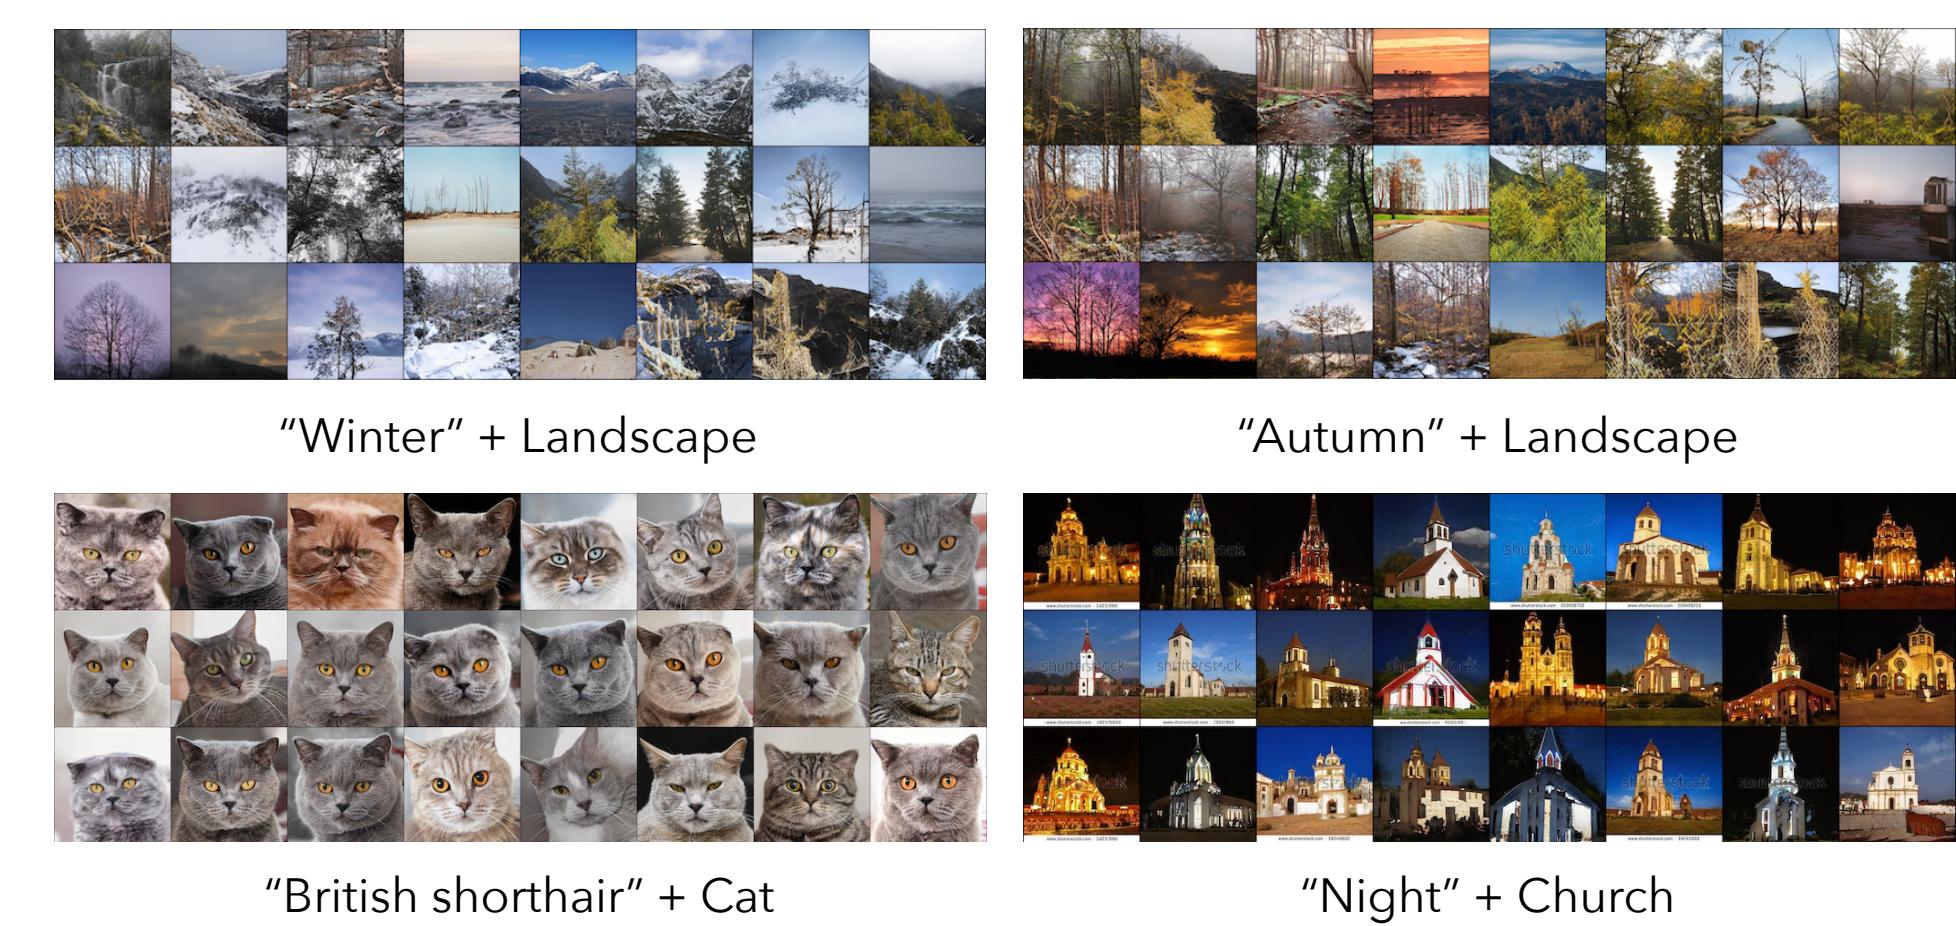
What do babies look like in StyleNeRF and Diffusion Autoencoder?



PromptGen for StyleNeRF

PromptGen for Diffusion AE

More domains

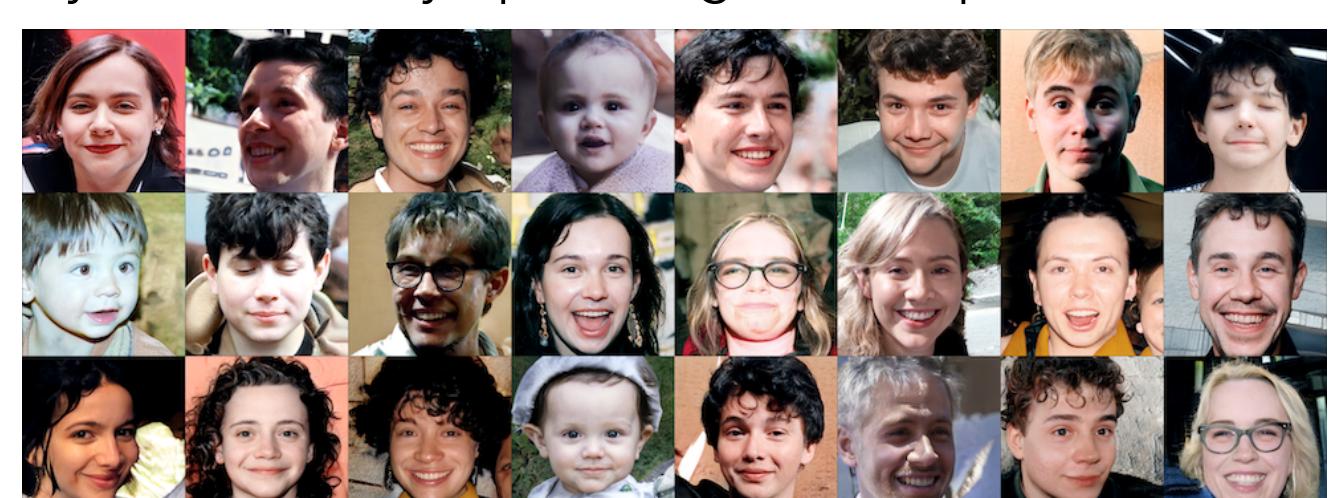


Comparison with ...

StyleGAN-NADA (fine-tuning model parameters)



StyleCLIP (locally optimizing latent representation)



Their failures show that *locally* optimizing model parameters or latent representation is not sufficient for distributional control.